

Shopify Intern Challenge

January 18, 2022

0.1 Shopify 2022 Data Science Intern Challenge

0.1.1 *By Anjali Chauhan*

Please complete the following questions, and provide your thought process/work. You can attach your work in a text file, link, etc. on the application page. Please ensure answers are easily visible for reviewers!

0.1.2 Question 1:

Given some sample data, write a program to answer the following: [click here](#) to access the required data set

On Shopify, we have exactly 100 sneaker shops, and each of these shops sells only one model of shoe. We want to do some analysis of the average order value (AOV). When we look at orders data over a 30 day window, we naively calculate an AOV of \$3145.13. Given that we know these shops are selling sneakers, a relatively affordable item, something seems wrong with our analysis.

- (i) Think about what could be going wrong with our calculation. Think about a better way to evaluate this data.
- (ii) What metric would you report for this dataset?
- (iii) What is its value?

0.1.3 Import libraries

```
[2]: import pandas as pd
import numpy as np
```

0.1.4 Reading the data

```
[3]: df = pd.read_csv('2019 Winter Data Science Intern Challenge Data Set - Sheet1.
↳csv')
df.head()
```

```
[3]:   order_id  shop_id  user_id  order_amount  total_items  payment_method \
0         1         53       746           224             2             cash
1         2         92       925            90             1             cash
2         3         44       861           144             1             cash
3         4         18       935           156             1      credit_card
```

4	5	18	883	156	1	credit_card
---	---	----	-----	-----	---	-------------


```

        created_at
0  2017-03-13 12:36:56
1  2017-03-03 17:38:52
2   2017-03-14 4:23:56
3  2017-03-26 12:43:37
4   2017-03-01 4:35:11

```

0.1.5 Exploratory Data Analysis

```
[4]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  -
0   order_id        5000 non-null   int64
1   shop_id         5000 non-null   int64
2   user_id         5000 non-null   int64
3   order_amount    5000 non-null   int64
4   total_items     5000 non-null   int64
5   payment_method  5000 non-null   object
6   created_at      5000 non-null   object
dtypes: int64(5), object(2)
memory usage: 273.6+ KB

```

```
[5]: df.describe()
```

```

[5]:
      order_id  shop_id  user_id  order_amount  total_items
count  5000.000000  5000.000000  5000.000000    5000.000000    5000.000000
mean    2500.500000    50.078800   849.092400    3145.128000     8.787200
std    1443.520003    29.006118    87.798982   41282.539349    116.320320
min         1.000000     1.000000   607.000000     90.000000     1.000000
25%    1250.750000    24.000000   775.000000    163.000000     1.000000
50%    2500.500000    50.000000   849.000000    284.000000     2.000000
75%    3750.250000    75.000000   925.000000    390.000000     3.000000
max    5000.000000   100.000000   999.000000  704000.000000   2000.000000

```

```

[6]: # Check for null values in the dataframe
null_values = df.isnull().sum().to_frame()
null_values.columns = ['null_values']
null_values

```

```

[6]:
      order_id  null_values
order_id      0

```

```
shop_id      0
user_id      0
order_amount  0
total_items  0
payment_method 0
created_at   0
```

```
[7]: # AOV calculated in the question
df.order_amount.sum()/len(df.order_amount) # mean
```

```
[7]: 3145.128
```

```
[8]: df.order_amount.median()
```

```
[8]: 284.0
```

```
[9]: # Correct Mean
round(df.order_amount.sum()/sum(df.total_items),2) # mean
```

```
[9]: 357.92
```

```
[13]: # Mode
df.order_amount.mode()
```

```
[13]: 0    153
dtype: int64
```

```
[10]: # Converting id's from integer type to string
df.order_id = df.order_id.astype(str)
df.shop_id = df.shop_id.astype(str)
df.user_id = df.user_id.astype(str)
```

```
[12]: df['item_price'] = (df.order_amount/df.total_items).astype(int)
df
```

```
[12]:
```

	order_id	shop_id	user_id	order_amount	total_items	payment_method	\
0	1	53	746	224	2	cash	
1	2	92	925	90	1	cash	
2	3	44	861	144	1	cash	
3	4	18	935	156	1	credit_card	
4	5	18	883	156	1	credit_card	
...	
4995	4996	73	993	330	2	debit	
4996	4997	48	789	234	2	cash	
4997	4998	56	867	351	3	cash	
4998	4999	60	825	354	2	credit_card	
4999	5000	44	734	288	2	debit	

	created_at	item_price
0	2017-03-13 12:36:56	112
1	2017-03-03 17:38:52	90
2	2017-03-14 4:23:56	144
3	2017-03-26 12:43:37	156
4	2017-03-01 4:35:11	156
...
4995	2017-03-30 13:47:17	165
4996	2017-03-16 20:36:16	117
4997	2017-03-19 5:42:42	117
4998	2017-03-16 14:51:18	177
4999	2017-03-18 15:48:18	144

[5000 rows x 8 columns]

```
[19]: # Average order amount for each shop (lowest 5)
df.groupby('shop_id')[['order_amount']].mean().sort_values('order_amount').
    ↪head()
```

```
[19]:      order_amount
shop_id
92      162.857143
2       174.327273
32      189.976190
100     213.675000
53      214.117647
```

```
[20]: # Average order amount for each shop (highest 5)
df.groupby('shop_id')[['order_amount']].mean().sort_values('order_amount').
    ↪tail()
```

```
[20]:      order_amount
shop_id
38      390.857143
90      403.224490
50      403.545455
78     49213.043478
42     235101.490196
```

```
[24]: # Shop 42
df[df.shop_id == '42')[['order_amount', 'total_items', 'item_price']].
    ↪value_counts()
```

```
[24]: order_amount  total_items  item_price
704000         2000         352          17
352             1         352          15
```

704	2	352	13
1056	3	352	3
1408	4	352	2
1760	5	352	1

dtype: int64

```
[25]: # Shop 78
df[df.shop_id == '78'][['order_amount', 'total_items', 'item_price']].
    value_counts()
```

```
[25]: order_amount  total_items  item_price
25725             1           25725         19
51450             2           25725         16
77175             3           25725          9
102900            4           25725          1
154350            6           25725          1
dtype: int64
```

0.1.6 Question 2

For this question you'll need to use SQL. Follow this link to access the data set required for the challenge. Please use queries to answer the following questions. Paste your queries along with your final numerical answers below.

(i) How many orders were shipped by Speedy Express in total?

Answer: 54 orders were shipped by Speedy Express in total

```
SELECT o.OrderID, o.ShipperID, s.ShipperID, s.ShipperName, COUNT(*) FROM [Orders] o
LEFT JOIN [Shippers] s ON o.ShipperID = s.ShipperID
WHERE s.ShipperName = 'Speedy Express'
```

(ii) What is the last name of the employee with the most orders?

Answer: Peacock is the last name of the employee with the most orders, about 40

```
SELECT e.LastName, COUNT(DISTINCT(o.OrderID)) AS total_orders FROM [Employees] e
LEFT JOIN [Orders] o
ON e.EmployeeID = o.EmployeeID
GROUP BY e.LastName
ORDER BY total_orders DESC
LIMIT 1
```

(iii) What product was ordered the most by customers in Germany?

Answer: Boston Crab Meat was the product ordered the most by customers in Germany, about 160. In this problem, product most ordered is defined as total quantity of the product ordered and not how many times it was ordered.

```
SELECT d.OrderID, p.ProductID, p.ProductName, SUM(d.Quantity) AS TotalQuantity FROM [Customers] c
LEFT JOIN [Orders] o
ON c.CustomerID = o.CustomerID
```

```
LEFT JOIN [OrderDetails] d
ON o.OrderID = d.OrderID
LEFT JOIN [PRODUCTS] p
ON d.ProductID = p.ProductID
WHERE c.Country = 'Germany'
GROUP BY p.ProductName
ORDER BY TotalQuantity DESC
LIMIT 1
```